

UNIVERSITY OF REGINA

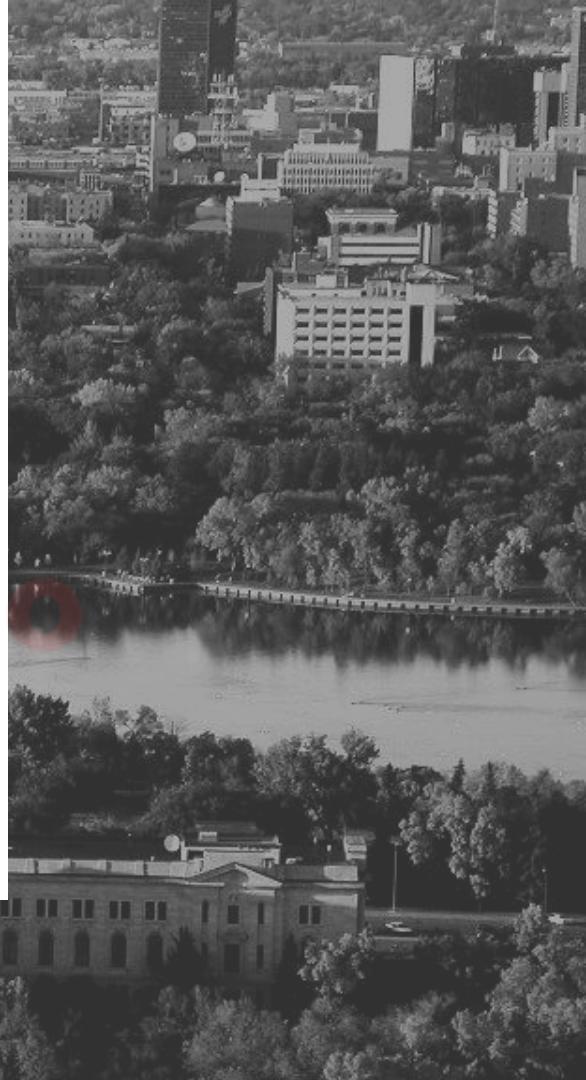
# CS330-001 INTRODUCTION TO OPERATING SYSTEMS

**andreeds.github.io**

ANDRÉ E. DOS SANTOS

[dossantos@cs.uregina.ca](mailto:dossantos@cs.uregina.ca)

[andreeds.github.io](http://andreeds.github.io)



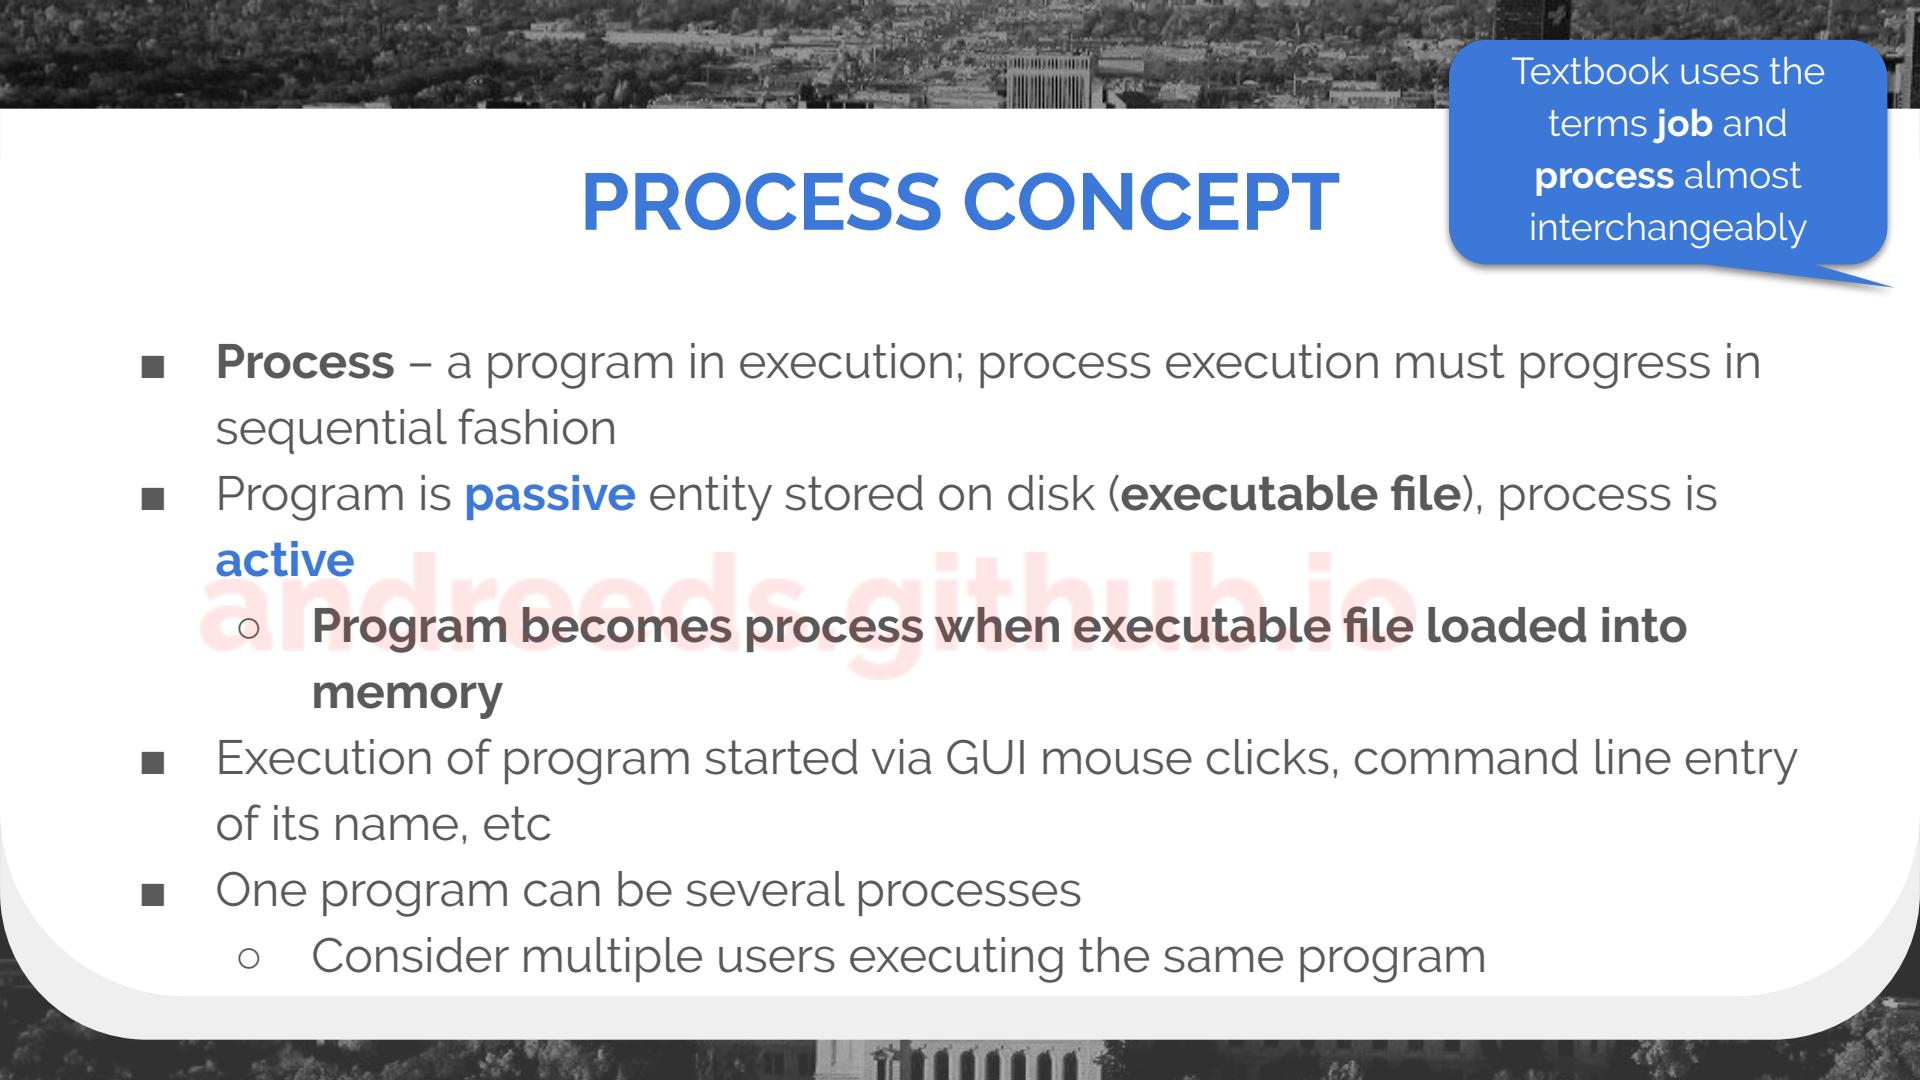
CS330-001  
INTRODUCTION TO  
OPERATING SYSTEMS

# PROCESSES

**andreeds.github.io**

ANDRÉ E. DOS SANTOS  
[dossantos@cs.uregina.ca](mailto:dossantos@cs.uregina.ca)  
[andreeds.github.io](http://andreeds.github.io)





Textbook uses the terms **job** and **process** almost interchangeably

# PROCESS CONCEPT

- **Process** – a program in execution; process execution must progress in sequential fashion
- Program is **passive** entity stored on disk (**executable file**), process is **active**
  - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
  - Consider multiple users executing the same program

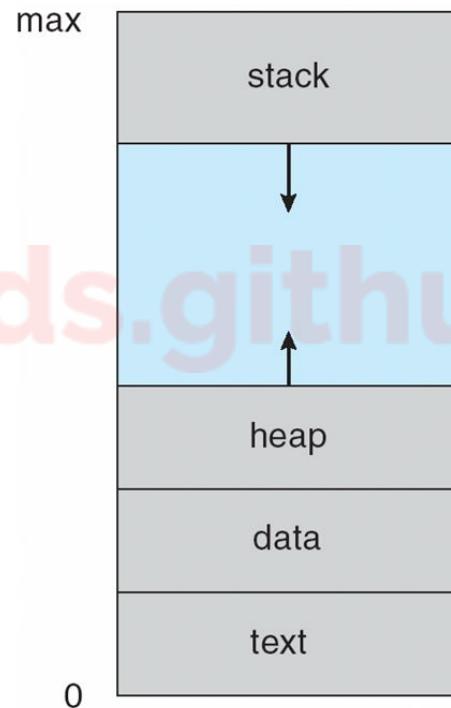
# PROCESS CONCEPT

- Multiple parts
  - The program code, also called **text section**
  - Current activity including **program counter**, processor registers
  - **Stack** containing temporary data
    - Function parameters, return addresses, local variables
  - **Data section** containing global variables
  - **Heap** containing memory dynamically allocated during run time

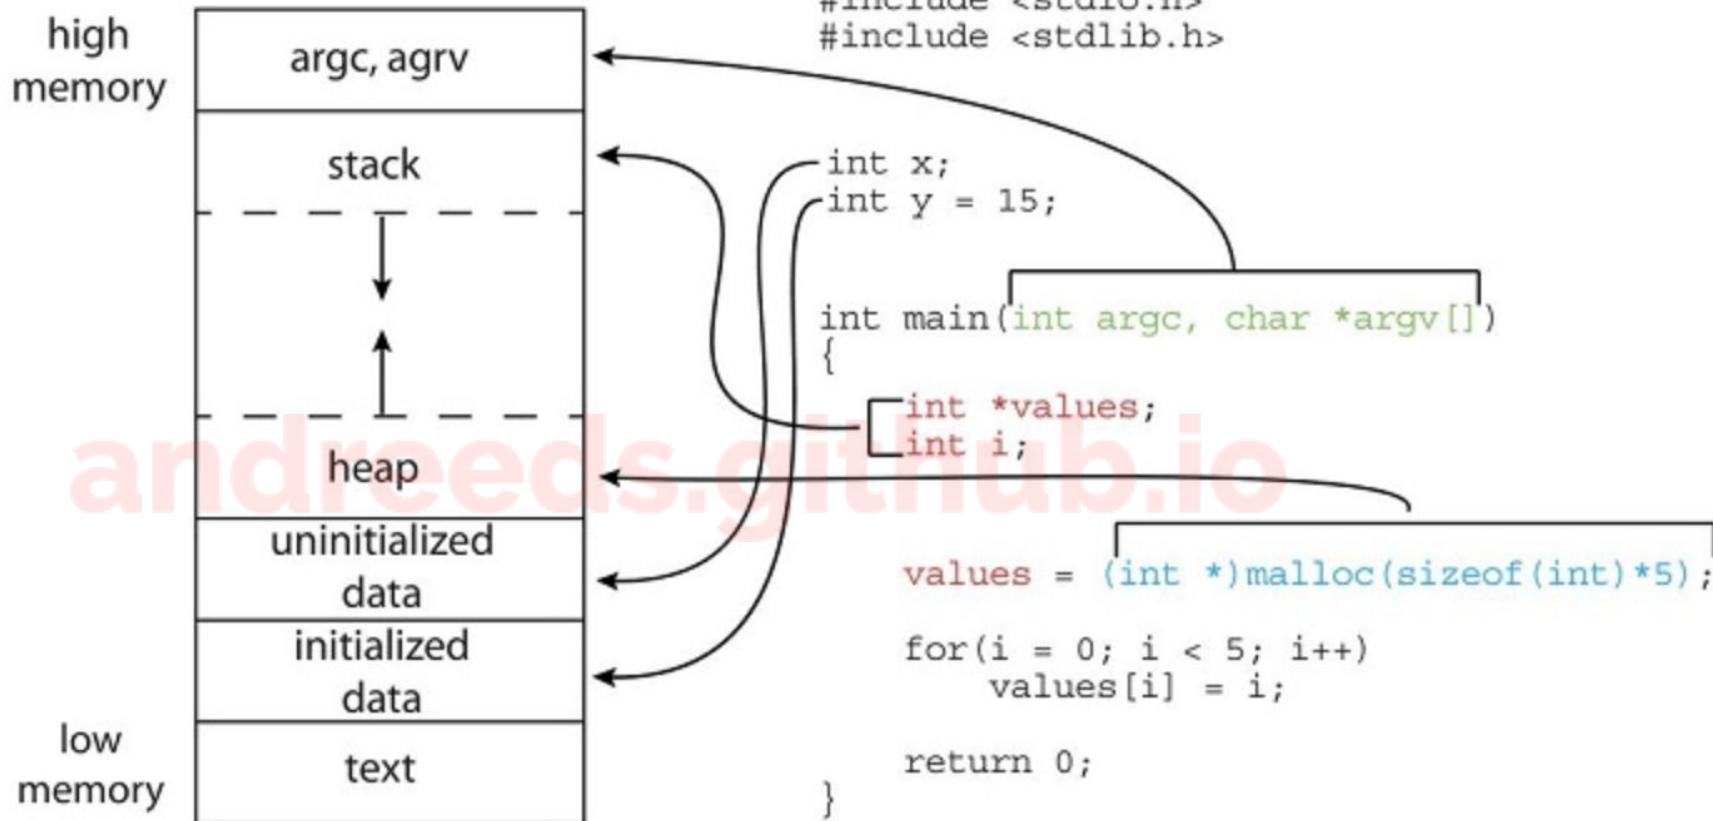
# THE PROCESS

p106

Layout of a process in memory

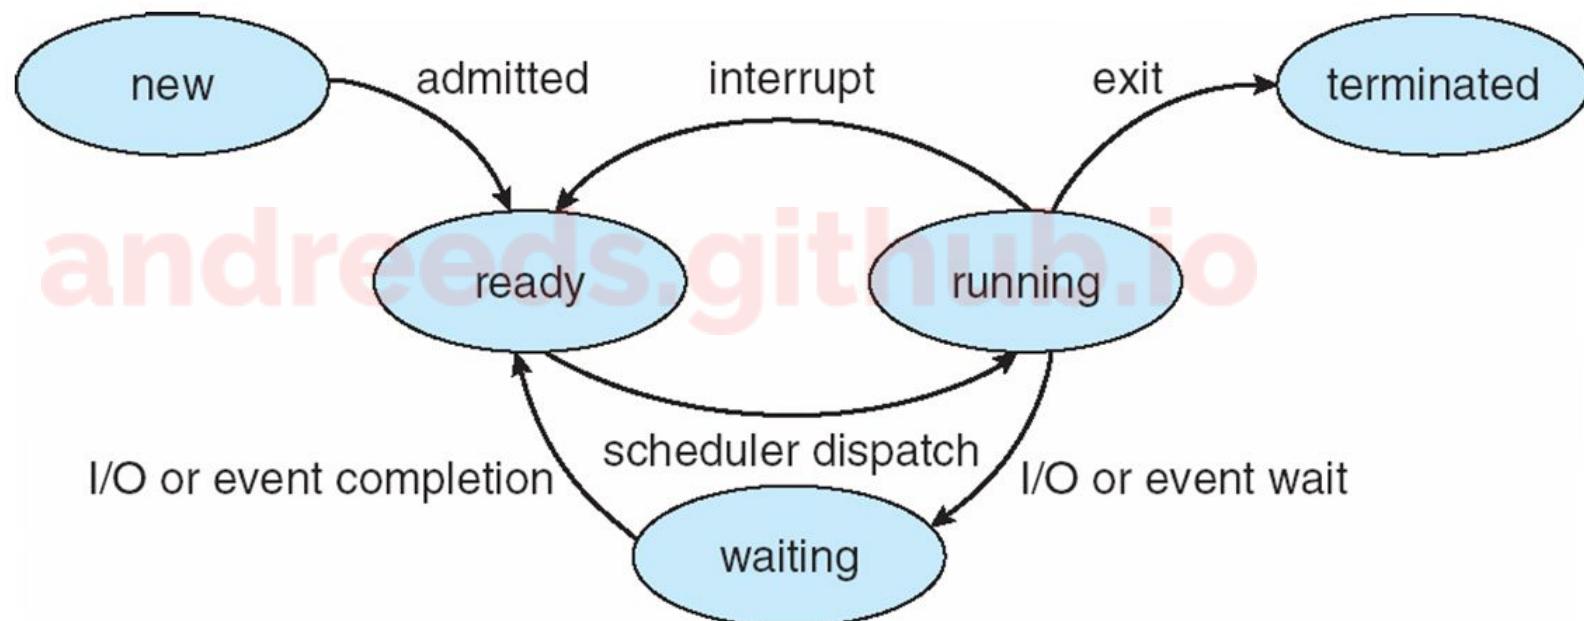


andreeeds.github.io



# PROCESS CONCEPT

## PROCESS STATE



# PROCESS CONCEPT

## PROCESS STATE

- Multiple parts
  - The program code, also called **text section**
  - Current activity including **program counter**, processor registers
  - **Stack** containing temporary data
    - Function parameters, return addresses, local variables
  - **Data section** containing **global** variables
  - **Heap** containing memory dynamically allocated during run time

# PROCESS CONCEPT

## PROCESS CONTROL BLOCK (PCB)

- Each process is **represented** in the operating system by a **process control block (PCB)**—also called a **task control block**.
- It contains many pieces of **information** associated with a specific process, including these:
  - Process state
  - Program counter
  - CPU registers
  - CPU scheduling information
  - Memory-management information
  - Accounting information
  - I/O status information

# PROCESS CONTROL BLOCK (PCB)



# PROCESS CONCEPT

## THREADS

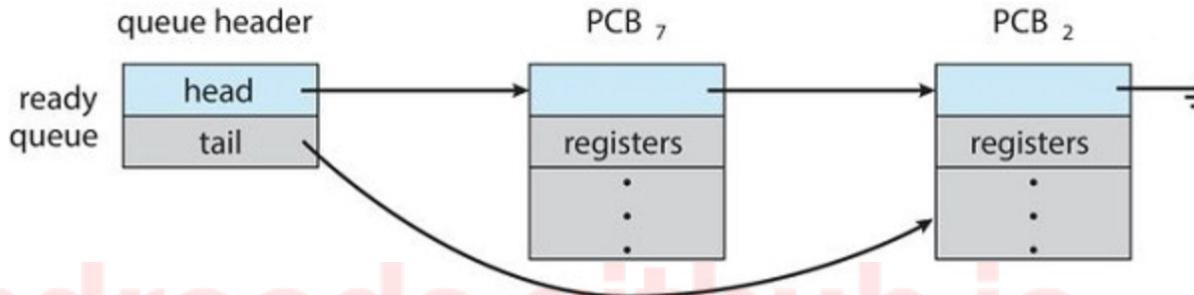
- So far, process has a single thread of execution
- Consider having multiple program counters per process
  - Multiple locations can execute at once
    - Multiple threads of control → **threads**
- Must then have storage for thread details, multiple program counters in PCB

# PROCESS SCHEDULING

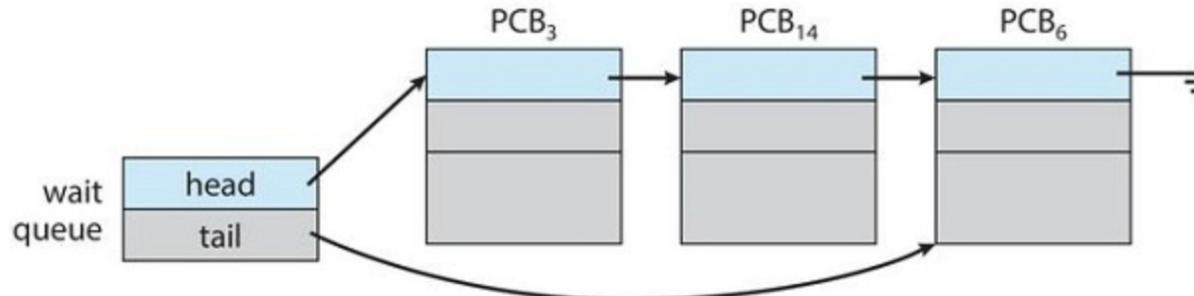
- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
  - **Job queue** – set of all processes in the system
  - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
  - **Device queues** – set of processes waiting for an I/O device
  - Processes migrate among the various queues

# READY QUEUE AND VARIOUS I/O DEVICE QUEUES

p109

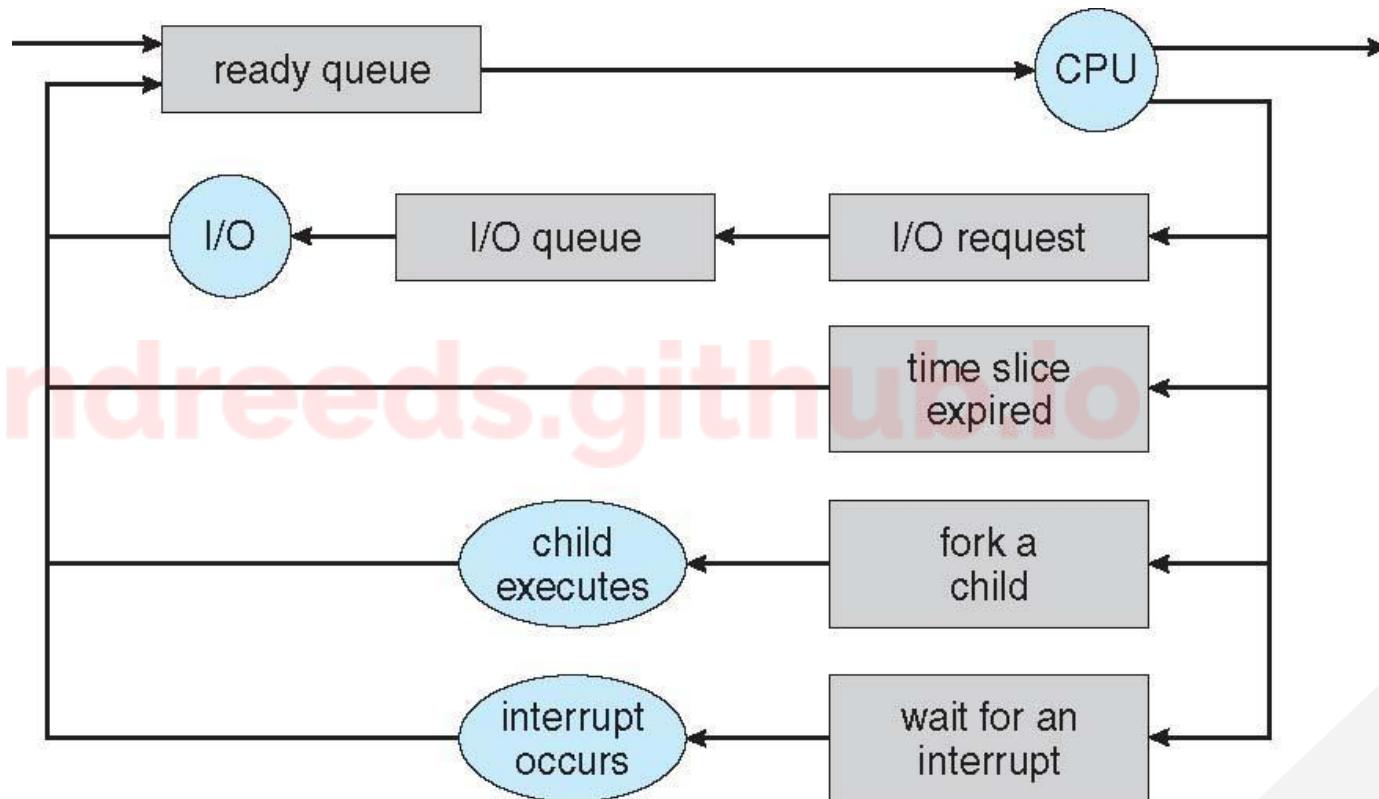


[andreevs.github.io](https://andreevs.github.io)



# QUEUEING DIAGRAM

p113



# PROCESS SCHEDULING

## CPU SCHEDULING

- A process migrates among the ready queue and various wait queues throughout its lifetime
- The role of the CPU scheduler is to select from among the processes that are in the ready queue and allocate a CPU core to one of them
- The CPU scheduler must select a new process for the CPU frequently

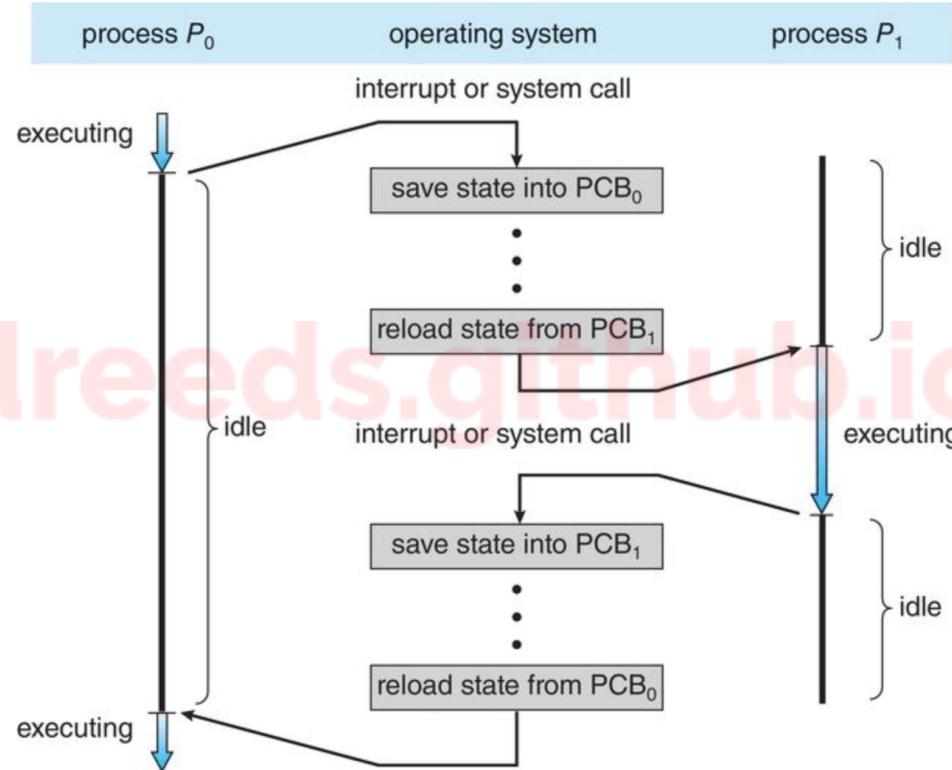
# PROCESS SCHEDULING

## CONTEXT SWITCH

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead;
  - the system does no useful work while switching
    - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support

# CONTEXT SWITCH

p114



# OPERATIONS ON PROCESSES

- System must provide mechanisms for:
  1. **process creation,**
  2. **process termination,**

and so on as detailed next

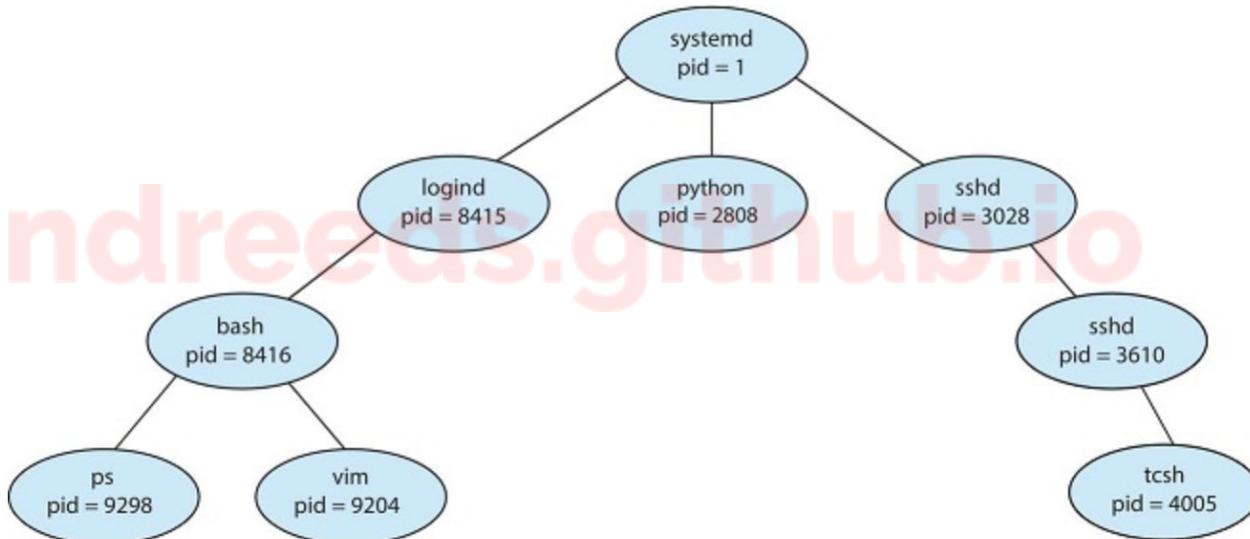
[andreeeds.github.io](https://andreeeds.github.io)

# OPERATIONS ON PROCESSES

## 1 PROCESS CREATION

- **Parent** process create **children** processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution options
  - Parent and children execute concurrently
  - Parent waits until children terminate

# A TREE OF PROCESSES IN LINUX



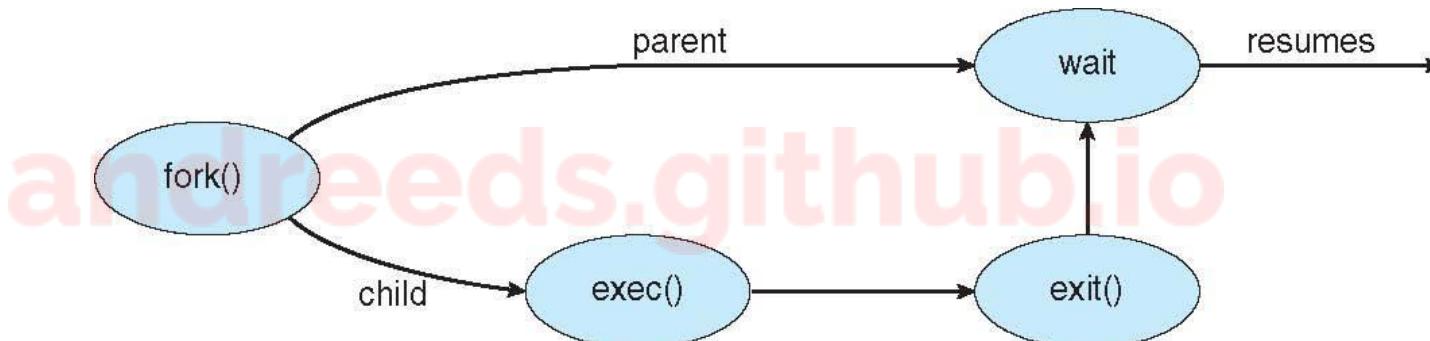
# OPERATIONS ON PROCESSES

## 1 PROCESS CREATION

- Address space
  - Child duplicate of parent
  - Child has a program loaded into it
- UNIX examples
  - **fork ()** system call creates new process
  - **exec ()** system call used after a **fork ()** to replace the process' memory space with a new program

# Process creation using the `fork()` system call

p119



# OPERATIONS ON PROCESSES

## 2 PROCESS TERMINATION

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
  - Returns status data from child to parent (via **wait()**)
  - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

# OPERATIONS ON PROCESSES

## 2 PROCESS TERMINATION

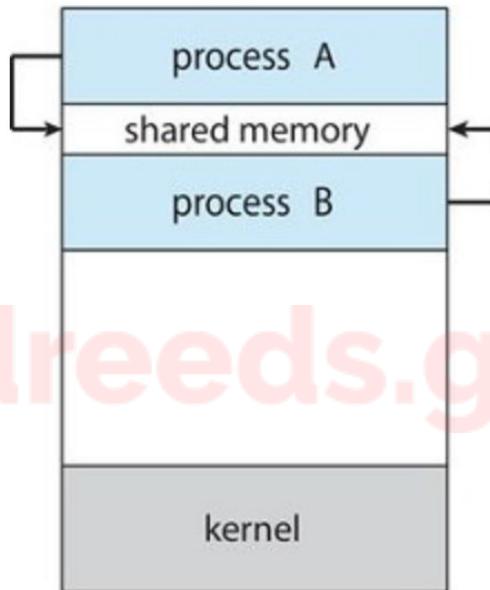
- Some operating systems do not allow child to exists if its parent has terminated.
- If a process terminates, then all its children must also be terminated.
  - **cascading termination.**
  - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process
  - **pid = wait(&status) ;**
- If no parent waiting (did not invoke **wait()**) process is a **zombie**
- If parent terminated without invoking **wait()**, process is an **orphan**

# INTERPROCESS COMMUNICATION

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
  - **Shared memory**
  - **Message passing**

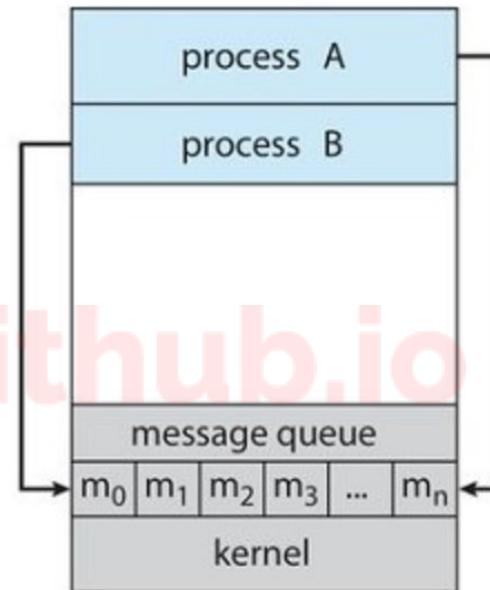
# COMMUNICATIONS MODELS

p125



(a)

**Shared memory**



(b)

**Message passing**

# PRODUCER-CONSUMER PROBLEM

Paradigm for cooperating processes,  
**producer process** produces information that is  
consumed by a **consumer process**

[andreevs.github.io](https://andreevs.github.io)

- **unbounded-buffer** places no practical limit on the size of the buffer
- **bounded-buffer** assumes that there is a fixed buffer size

# INTERPROCESS COMMUNICATION

## SHARED MEMORY

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory

# INTERPROCESS COMMUNICATION

## MESSAGE PASSING

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send (message)**
  - **receive (message)**
- The **message** size is either fixed or variable

# INTERPROCESS COMMUNICATION

## MESSAGE PASSING

- Implementation of communication link

- Physical:

- Shared memory
    - Hardware bus
    - Network

- Logical:

- **Direct or indirect**
    - **Synchronous or asynchronous**
    - **Automatic or explicit buffering**

andreeds.github.io

# INTERPROCESS COMMUNICATION

## DIRECT COMMUNICATION

- Processes must name each other explicitly:
  - **send (P, message)** – send a message to process **P**
  - **receive (Q, message)** – receive a message from process **Q**
- Properties of **communication link**
  - **Links** are established automatically
  - A link is associated with exactly one pair of communicating processes
  - **Between each pair there exists exactly one link**
  - The link may be unidirectional, but is usually bi-directional

# INTERPROCESS COMMUNICATION

## INDIRECT COMMUNICATION

- Messages are directed and received from mailboxes (also referred to as ports)
  - **Each mailbox has a unique id**
  - Processes can communicate only if they share a mailbox
- Properties of communication link
  - Link established only if processes share a common mailbox
  - **A link may be associated with many processes**
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional

# INTERPROCESS COMMUNICATION

## INDIRECT COMMUNICATION

- Operations
  - create a new mailbox (port)
  - send and receive messages through mailbox
  - destroy a mailbox
- Primitives are defined as:
  - **send (A, message)** – send a message to mailbox A
  - **receive (A, message)** – receive a message from mailbox A

# INTERPROCESS COMMUNICATION

## INDIRECT COMMUNICATION

- Mailbox sharing
  - P<sub>1</sub>, P<sub>2</sub>, and P<sub>3</sub> share **mailbox A**
  - P<sub>1</sub>, sends; P<sub>2</sub> and P<sub>3</sub> receive
  - Who gets the message?
- Solutions
  - Allow a link to be associated with at most two processes
  - Allow only one process at a time to execute a receive operation
  - Allow the system to select arbitrarily the receiver
    - Sender is notified who the receiver was

# INTERPROCESS COMMUNICATION

## SYNCHRONIZATION

- Message passing may be either **blocking** or **non-blocking**
- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking send** -- the sender sends the message and continue
  - **Non-blocking receive** -- the receiver receives:
    - A valid message, or
    - Null message
- Different combinations possible
- If both send and receive are blocking, we have a **rendezvous**

# INTERPROCESS COMMUNICATION

## BUFFERING

- Buffers are generally used when there is a difference between the rate at which data is received and the rate at which it can be processed.
- Queue of messages attached to the **link**
- implemented in one of three ways
  1. Zero capacity – no messages are queued on a link
    - Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of n messages
    - Sender must wait if link full
  3. Unbounded capacity – infinite length
    - Sender never waits

# **COMMUNICATIONS IN CLIENT-SERVER SYSTEMS**

- Sockets
- Remote Procedure Calls
- Pipes

**ahdreeeds.github.io**

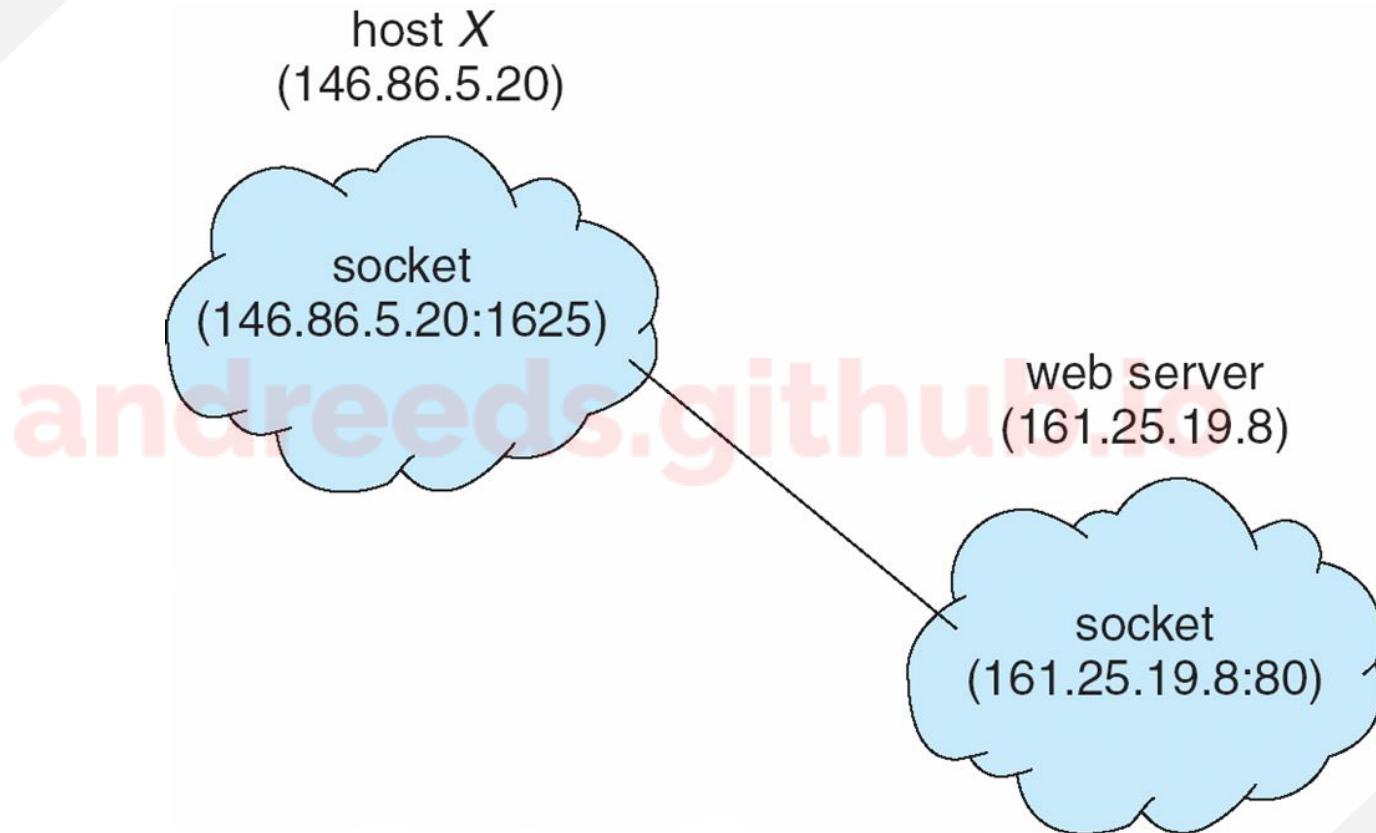
# COMMUNICATIONS IN CLIENT-SERVER SYSTEMS

## SOCKETS

- A **socket** is defined as an endpoint for communication
- Concatenation of **IP address** and **port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below **1024** are **well known**, used for standard services
- Special IP address **127.0.0.1 (loopback)** to refer to system on which process is running

# SOCKET COMMUNICATION

p147



# COMMUNICATIONS IN CLIENT-SERVER SYSTEMS

## REMOTE PROCEDURE CALLS

- **Remote procedure call (RPC)** is a protocol that one program can use to request a service from a program located in another computer on a network without having to understand the network's details

[andreevs.github.io](https://andreevs.github.io)

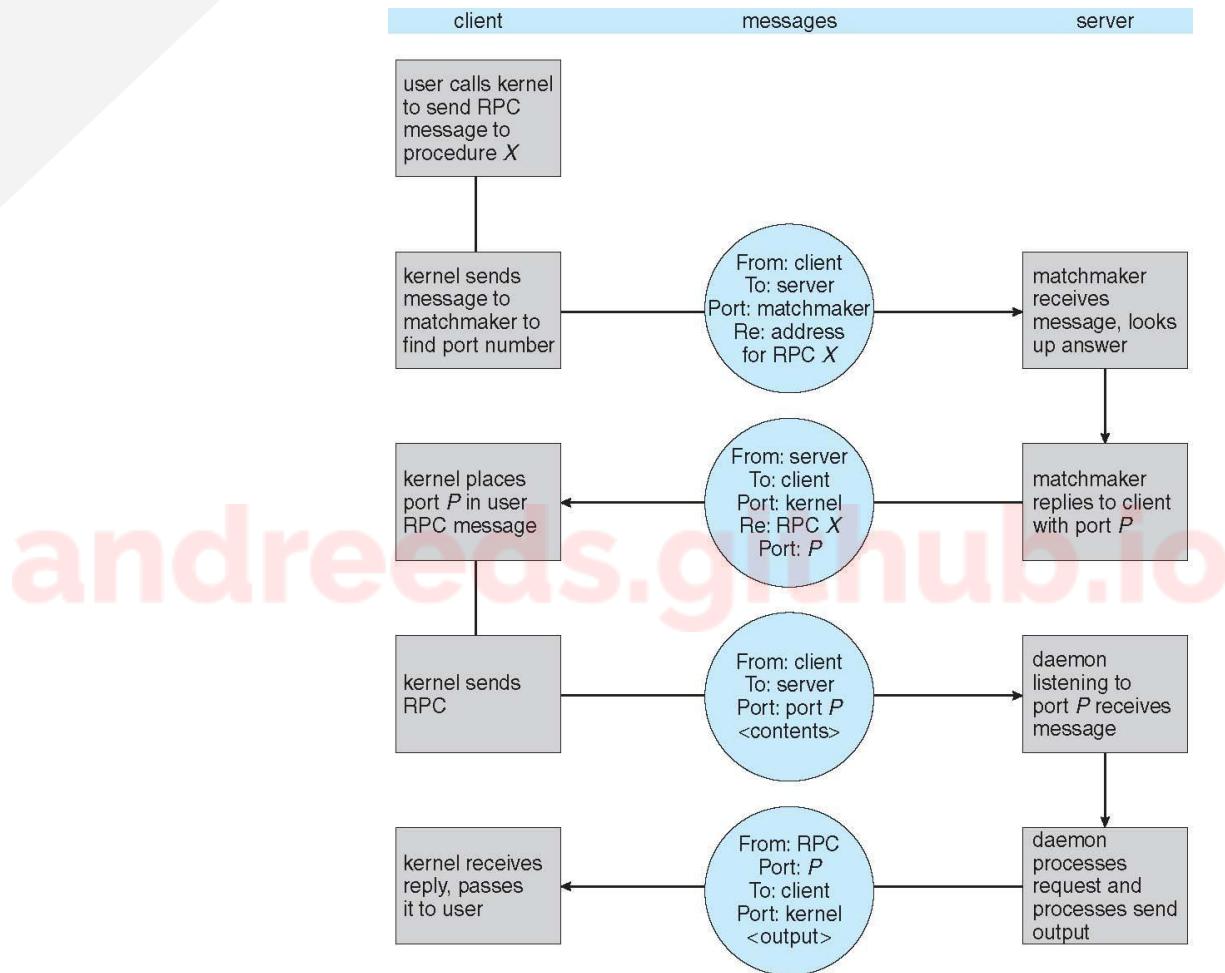
# COMMUNICATIONS IN CLIENT-SERVER SYSTEMS

## REMOTE PROCEDURE CALLS

- **Remote procedure call (RPC)** is a protocol that one program can use to request a service from a program located in another computer on a network without having to understand the network's details
- When program statements that use RPC framework are compiled into an executable program, a **stub** is included in the compiled code that acts as the representative of the remote procedure code
- When the program is run and the procedure call is issued, the stub receives the request and forwards it to a client runtime program in the local computer

**stub**

a small program routine that substitutes for a longer program, possibly to be loaded later or that is located remotely



# COMMUNICATIONS IN CLIENT-SERVER SYSTEMS

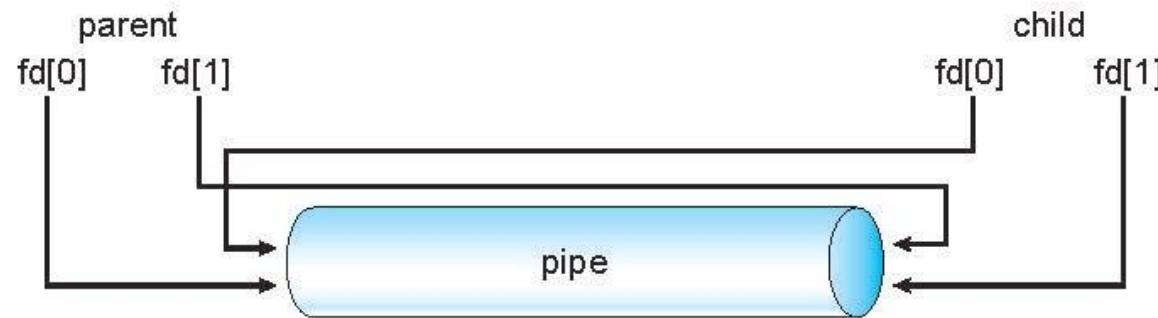
## PIPES

- Acts as a conduit allowing two processes to communicate
- Issues:
  - Is communication unidirectional or bidirectional?
  - In the case of two-way communication, is it half or full-duplex?
  - Must there exist a relationship (i.e., **parent-child**) between the communicating processes?
  - Can the pipes be used over a network?
- **Ordinary pipes** – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- **Named pipes** – can be accessed without a parent-child relationship.

# COMMUNICATIONS IN CLIENT-SERVER SYSTEMS

## ORDINARY PIPES

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



# COMMUNICATIONS IN CLIENT-SERVER SYSTEMS

## NAMED PIPES

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems



# REVIEW QUESTIONS

## PROCESSES

The \_\_\_\_ of a process contains temporary data such as function parameters, return addresses, and local variables.

- A. text section
- B. data section
- C. program counter
- D. stack

A process may transition to the Ready state by which of the following actions?

- A. Completion of an I/O event
- B. Awaiting its turn on the CPU
- C. Newly-admitted process
- D. All of the above

A process control block \_\_\_\_.

- A. includes information on the process's state
- B. stores the address of the next instruction to be processed by a different process
- C. determines which process is to be executed next
- D. is an example of a process queue



# REVIEW QUESTIONS PROCESSES

A \_\_\_\_\_ saves the state of the currently running process and restores the state of the next process to run.

- A. save-and-restore
- B. state switch
- C. context switch
- D. none of the above

When a child process is created, which of the following is a possibility in terms of the execution or address space of the child process?

- A. The child process runs concurrently with the parent.
- B. The child process has a new program loaded into it.
- C. The child is a duplicate of the parent.
- D. All of the above



# REVIEW QUESTIONS PROCESSES

→ Which of the following statements is true?

- A. Shared memory is typically faster than message passing.
- B. Message passing is typically faster than shared memory.
- C. Message passing is most useful for exchanging large amounts of data.
- D. Shared memory is far more common in operating systems than message passing.

→ A blocking send() and blocking receive() is known as a(n) \_\_\_\_\_

- A) synchronized message
- B) rendezvous
- C) blocked message
- D) asynchronous message

→ "Under indirect communication, each process that wants to communicate must explicitly name the recipient or sender of the communication." True or False?



# REVIEW QUESTIONS PROCESSES

Imagine that a host with IP address 150.55.66.77 wishes to

download a file from the web server at IP address  
202.28.15.123. Select a valid socket pair for a connection  
between this pair of hosts.

- A) 150.55.66.77:80 and 202.28.15.123:80
- B) 150.55.66.77:150 and 202.28.15.123:80
- C) 150.55.66.77:2000 and 202.28.15.123:80
- D) 150.55.66.77:80 and 202.28.15.123:3500

When a process creates a new process using the  
`fork()` operation, which of the following states is  
shared between the parent process and the child  
process?

- a. Stack
- b. Heap
- c. Shared memory segments

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int value = 5;

int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) /* parent process */

        wait(NULL);
        printf("PARENT: value = %d",value); /* LINE A */

        return 0;
}
```

Explain what the output will be  
at LINE A.

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    /* fork a child process */
    fork();

    /* fork another child process */
    fork();

    /* and fork another */
    fork();

    return 0;
}
```

Including the initial parent process, how many processes are created by the program?