

UNIVERSITY OF REGINA

CS330-001  
INTRODUCTION TO  
OPERATING  
SYSTEMS  
**andreeds.github.io**

ANDRÉ E. DOS SANTOS  
**dossantos@cs.uregina.ca**  
**andreeds.github.io**



CS330-001  
INTRODUCTION TO  
OPERATING SYSTEMS

# PROGRAMMING PROCESSES

C

[andreeds.github.io](https://andreeds.github.io)

ANDRÉ E. DOS SANTOS  
[dossantos@cs.uregina.ca](mailto:dossantos@cs.uregina.ca)  
[andreeds.github.io](https://andreeds.github.io)



# UNIX/LINUX PROCESSES

[andreevs.github.io](https://andreevs.github.io)

# THE `fork` FUNCTION

- A process can create a new process by calling the `fork` function

```
#include <unistd.h>
pid_t fork (void);
```

- The calling process becomes the *parent*
- The new process is called the *child*
- The parent and child have different process IDs
- The `pid_t` type is an unsigned integer
- In the parent, the call to `fork` returns the process ID of the child process if the creation of the child process was successful. Otherwise, it returns `-1`
- In the child, the call to `fork` always returns `0`.
- The parent and child have their own separate PCBs, so do not share any data through like-named variables
- Both the parent and child continue at the instruction following the call to `fork`

# THE getpid AND getppid FUNCTION

- The `getpid` and `getppid` functions return the process ID and parent process ID, respectively

```
#include <unistd.h>
pid_t getpid (void);
pid_t getppid (void);
```

- Neither `getpid` nor `getppid` can return an error.

# UNIX/LINUX PROCESSES

simpleFork.c

URCourses

```
#include <stdio.h>
#include <unistd.h>

int main ()
{
    int x = 0;
    fork ();
    x = 1;
    printf ("I am process %d, ", getpid ());
    printf ("my parent is %d, ", getppid ());
    printf ("and my x is %d.\n", x);

    return 0;
}
```

# UNIX/LINUX PROCESSES

betterFork.c

URCourses

```
#include <stdio.h>
#include <unistd.h>

int main ()
{
    pid_t childPid = 0;
    int x = 0;

    childPid = fork ();
    if (childPid == -1)
    {
        fprintf (stderr, "Process %d failed to
fork!\n", getpid ());
        return 1;
    }
}
```

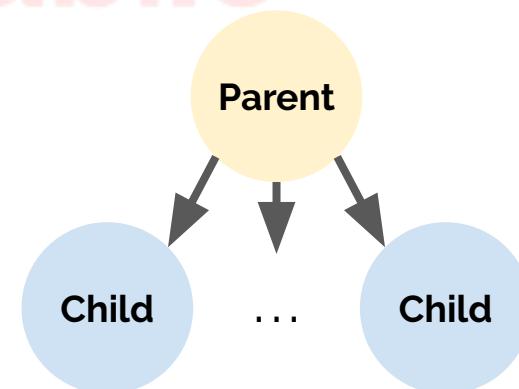
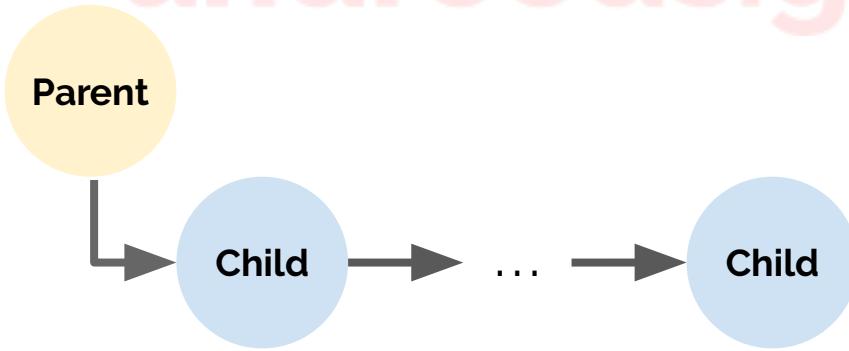
# CHAIN AND FAN

## Chain

```
pid_t childpid = 0;  
for (i=1;i<n;i++)  
    if (childpid = fork())  
        break;
```

## Fan

```
pid_t childpid = 0;  
for (i=1;i<n;i++)  
    if (childpid = fork() <= 0)  
        Break;
```



# UNIX/LINUX PROCESSES

processChain.c

URCourses

```
#include <stdio.h>
#include <unistd.h>

int main ()
{
    pid_t childPid = 0;
    int i;
    int n = 4;
    int x = 0;

    for (i = 1; i < n; i++)
    {
        if (childPid = fork ())

```

# UNIX/LINUX PROCESSES

processFan.c

URCourses

```
#include <stdio.h>
#include <unistd.h>

int main ()
{
    pid_t childPid = 0;
    int i;
    int n = 4;
    int x = 0;

    for (i = 1; i < n; i++)
    {
        x = i;
```



**YOU ARE ABSOLUTELY FORBIDDEN TO EXPERIMENT WITH  
PROGRAMS USING `fork()` ON HERCULES!!**

# WHY?

- When using and experimenting with **fork**, you may set up a process that creates some number of child processes, and then each of these processes creates some number of child processes, and then each of these processes...
  - you get the picture
- Now let's assume that you have a bug in your program so that the creation of child processes is not going to stop
  - i.e., the termination condition is never satisfied
- Even if each process creates just two child processes, after one generation, you'll have two more processes, after two generations, four more, after three, eight more, ..., after n,  $2^n$ 
  - clearly, the number of processes created is growing exponentially

# WHY?

- These processes are created so fast that you will be unable to kill them fast enough to prevent Hercules from becoming overwhelmed
- Essentially, as soon as some number of processes is killed, the newly available resources will quickly be overwhelmed with new processes
  - This is called a **fork** BOMB!!
- SO, WHEN USING AND EXPERIMENT WITH **fork**, DO IT ON THE LINUX MACHINES ONLY!
  - Beware, it may happen to YOU
  - A repeat performance could cause your computing privileges to be restricted.

# THE sleep FUNCTION

- `#include <unistd.h>`
- `unsigned sleep (unsigned seconds) ;`
  - The `sleep` function returns `0` after the specified time has elapsed.

[andreevs.github.io](http://andreevs.github.io)

# UNIX/LINUX PROCESSES

processSleepyFan.c

URCourses

```
#include <stdio.h>
#include <unistd.h>

int main ()
{
    pid_t childPid = 0;
    int i;
    int n = 4;
    int x = 0;

    for (i = 1; i < n; i++)
    {
        if ((childPid = fork ()) <= 0)
```

# THE wait FUNCTION

- The wait and waitpid functions cause the parent to block until the child has finished executing

```
#include <sys/wait.h>
pid_t wait (int *status);
pid_t waitpid (pid_t pid, int *status, int options);
```

- The **wait** function causes the parent to block until it receives a signal that **any** child has finished executing

# THE `wait` FUNCTION

- The `waitpid` function causes the parent to block until it receives a signal that a **particular** child has finished executing.
- The `status` parameter is a pointer to a location for returning the status.
- The `pid` parameter determines how the parent will respond.
  - If pid is -1, the parent waits for any child (i.e., `waitpid` behaves like `wait` if `options` is also set to 0).
  - If pid is greater than 0, the parent waits for the child whose process ID is pid.
- The `options` parameter determines how the parent will handle the call to the `waitpid` function
  - i.e., blocking or non-blocking

# THE `wait` FUNCTION

- The `wait` and `waitpid` functions return the process ID of the child whose status is being reported
- If an error occurs, they return `-1` and set `errno`

[andreevs.github.io](https://andreevs.github.io)

# THE `wait` FUNCTION

- The `wait` and `waitpid` functions return the process ID of the child whose status is being reported
- If an error occurs, they return `-1` and set `errno`

[andreevs.github.io](https://andreevs.github.io)

# UNIX/LINUX PROCESSES

patientParentOfOne.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main ()
{
    pid_t childPid;
    int status;
    pid_t waitPid;

    childPid = fork ();
    if (childPid == 0)
    {
```

# UNIX/LINUX PROCESSES

patientParentOfMany.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main ()
{
    pid_t childPid;
    int i;
    int n = 4;
    int status;
    pid_t waitPid;

    for (i = 1; i < n; i ++)
```

# ERROR HANDLING IN SYSTEM CALLS

- All system calls return **-1** if an error occurs
- There are a number of helpful facilities to
- diagnose the problem
  - the **errno** global variable

**andreeeds.github.io**

# ERROR HANDLING IN SYSTEM CALLS

- **errno**
  - A **global variable** called **errno** in every process
  - At process creation time, **errno** is initialized to **0**
  - When a system call error occurs, **errno** is set
  - See **/usr/include/sys/errno.h**
    - The **<errno.h>** header shall provide a declaration for **errno** and give positive values for the following symbolic constants
  - *Examples*
    - **EAGAIN** = "Cannot fork process: System Process Limit Reached"
    - **ENOMEM** = "Cannot fork process: Out of memory"



# COMMAND LINE ARGUMENTS

[andreeeds.github.io](https://andreeeds.github.io)

# ARGUMENTS TO `main()`

- It is possible to pass arguments to `main` from a command line by including parameters `argc` and `argv` in the parameter list of main
- `int main (int argc, char **argv)`  
or
- `int main (int argc, char *argv [])`
  - A command line consists of arguments separated by white space
  - An argument may not contain whitespace unless surrounded by double quotes.
  - The shell parses the command line arguments into tokens to create a parameter called `argv` consisting of an array of `argc` pointers to C-strings whose last entry is a `NULL` pointer

## Example

```
% message worlds best boss
```

<b>argv</b>	[0]	→	'm'	'e'	's'	's'	'a'	'g'	'e'	'\0'
	[1]	→	'w'	'o'	'r'	'l'	'd'	's'	'\0'	
	[2]	→	'b'	'e'	's'	't'	'\0'			
	[3]	→	'b'	'o'	's'	's'	'\0'			
	[4]	→	∅							

```
argc = 4
```

- The first argument, **message**, is the program name
- The arguments to **message** are **worlds**, **best**, and **boss**

# UNIX/LINUX PROCESSES

commandLine.c

URCourses

```
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char *argv [])
{
    int i = 0;
    double d;

    if (argc != 5)
    {
        fprintf (stderr, "Usage: cl arg1
arg2 arg3 9.9\n");
        return 1;
    }

    d = atof(argv[3]);
    if (d >= 9.9)
        i = 1;
    else
        i = 0;
}
```

# THE `exec` FUNCTION

- The `exec` family of function calls provides a facility for overlaying the process image (i.e., the PCB) of the calling process with a new image from a process that is **different from that of the parent**
- An `exec` function call is used in combination with a `fork` function call
  - The `fork` function call creates a new child process
  - The `exec` function call starts a new process that overlays the child process
    - i.e., the code for the child process that resulted from the `fork` will no longer exist
  - The new process executes its code
  - The parent process continues to execute its original code

# THE exec FUNCTION

```
#include <unistd.h>

int execl (const char *path, const char *arg0, ..., char *(0));
int execv (const char *path, char *const argv []);
```

# THE exec1 PARAMETERS

- If you know the number of command line arguments at compile time
- The **path** parameter is the pathname of an executable file
- The **file** parameter is the name of an executable file
- The **arg<sub>i</sub>** parameters are the individual command line arguments
- The last **arg<sub>i</sub>** parameter should be a **NULL** pointer
  - `exec1 ("/bin/ls", "ls", "-l", NULL);`

# UNIX/LINUX PROCESSES

forkExec1.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main ()
{
    pid_t childPid;
    int status;
    pid_t waitPid;

    childPid = fork ();
    if (childPid == 0)
    {
```

# THE execv PARAMETERS

- If the command line arguments are determined (i.e., built) at run time
- The `path` parameter is the pathname of an executable file
- The `file` parameter is the name of an executable file
- The `argv []` parameter is an argument array whose last element is `NULL`
  - `execv (argv [1], &argv [1]);`

# UNIX/LINUX PROCESSES

forkExecv.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main (int argc, char *argv [])
{
    pid_t childPid;
    int status;
    pid_t waitPid;

    childPid = fork ();
    if (childPid == 0)
    {
```

# THE execvp PARAMETERS

- If the command line arguments are determined (i.e., built) at run time\*
- It has different behavior depending on the value of the first argument.
  - If it contains a slash, then \*file is treated as a pathname and **execvp** will behave like **execv**
  - **execvp (" /bin/ls" , &argv [0] ) ;**
  - Or
  - **char \*myarg[3];**  
**myarg[0] = strdup("ls");**  
**myarg[1] = strdup("-l");**  
**myarg[2] = NULL;**  
**execvp (" /bin/ls" , &myarg [0] ) ;**

# THE execvp PARAMETERS

or

- ```
char *myarg[3];
myarg[0] = strdup("ls");
myarg[1] = strdup("-l");
myarg[2] = NULL;
execvp ("/bin/ls", &myarg [0]);
```

or

- ```
char *myarg[4];
myarg[0] = strdup("ls");
myarg[1] = strdup("-l");
myarg[2] = argv[1];
myarg[3] = NULL;
execvp ("/bin/ls", &myarg [0]);
```

# UNIX/LINUX PROCESSES

## forkExecvp\_1.c

URCourses

```
// Compile:  
//         % gcc forkExecvp_1.c -o parent  
// Call:  
//         % parent -l  
  
// If  
//         execvp ("/bin/ls", &argv [1]);  
// Then Call  
//         % parent ls -l  
  
#include <stdio.h>  
#include <unistd.h>  
#include <sys/wait.h>
```

# UNIX/LINUX PROCESSES

## forkExevp\_2.c

URCourses

```
// Compile:  
//         gcc forkExevp_2.c -o parent  
// Call:  
//         % parent  
  
#include <stdio.h>  
#include <unistd.h>  
#include <sys/wait.h>  
#include <string.h> // because of strdup  
  
int main (int argc, char *argv [])  
{
```

# UNIX/LINUX PROCESSES

## forkExevp\_3.c

URCourses

```
// Compile:  
//         gcc forkExevp_3.c -o parent  
// Call:  
//         % parent -r  
//////// -r for reversed order, you may  
try another command options like -s, -S,  
-i, --color  
  
#include <stdio.h>  
#include <unistd.h>  
#include <sys/wait.h>  
#include <string.h> // because of strdup
```