



UNIVERSITY OF REGINA

CS330-001
**INTRODUCTION TO
OPERATING
SYSTEMS**

andreeds.github.io

ANDRÉ E. **DOS SANTOS**
dossantos@cs.uregina.ca
andreeds.github.io

CS330-001
INTRODUCTION TO
OPERATING SYSTEMS

PROGRAMMING SOCKETS

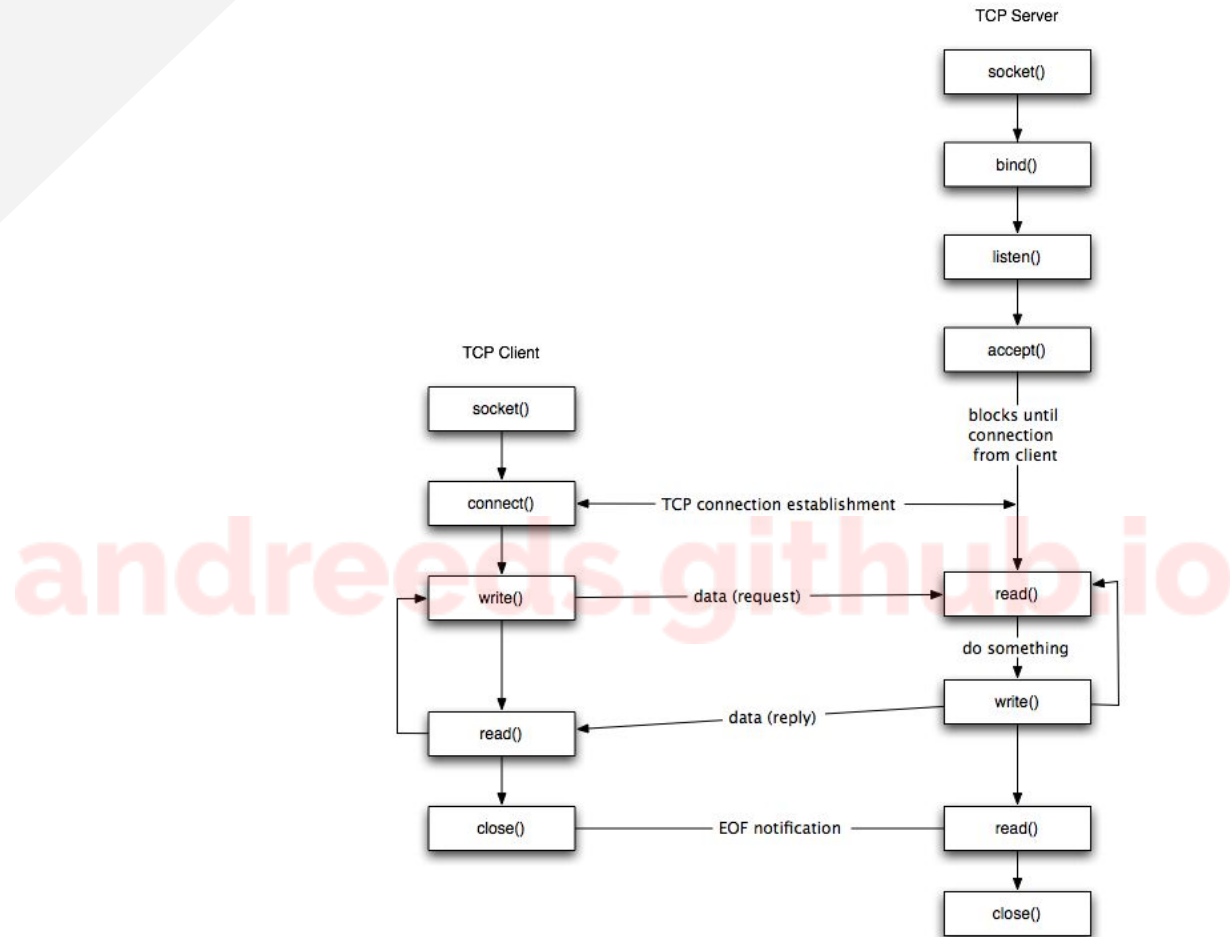
C

andreeds.github.io

ANDRÉ E. **DOS SANTOS**
dossantos@cs.uregina.ca
andreeds.github.io

PROGRAMMING SOCKETS

- BSD sockets
 - a fundamental component in implementing applications using the message passing paradigm on UNIX systems
- BSD sockets (API) available to programmers on UNIX systems
- A socket is an endpoint used by a process for bidirectional communication with another socket/process, usually across a network
- Generally, a socket is described by an IP address and a port number



SOCKET ADDRESSES

- Many of the networking system calls require a pointer to a socket address structure as an argument

```
#include <sys/types.h>

struct sockaddr
{
    u_short sa_family;    //address family
    char    sa_data [14]; //up to 14 bytes of protocol-specific address
};
```

SOCKET ADDRESSES

- For Internet protocol-specific addresses, an Internet address structure is required

```
#include <netinet/in.h>

struct in_addr
{
    u_long s_addr;    // 32 bit net id / host id
};

struct sockaddr_in
{
    short      sin_family;    // AF_INET
    u_short    sin_port;      // 16 bit port number
    struct in_addr sin_addr;   // 32 bit net id / host id
    char       sin_zero [8];  // unused
};
```

- The data type definitions (i.e., typedefs) for `u_short` and `u_long` are provided in `<sys/types.h>`

ELEMENTARY SOCKET SYSTEM CALLS

A **connection** describes the **association** (i.e., communication link) **between two processes**

- An **association** is a 5-tuple that completely specifies the two processes that make up a connection:

```
{protocol, local-address, local-process, foreign-address, foreign-process}
```

- The **local-address** and **foreign-address** specify the **network ID** and **host ID** of the **local** and **foreign hosts**, respectively
- The **local-address** and **foreign-process** identify the specific port on each end of the connection

- Example

```
{TCP, 192.43.235.2, 1500, 192.43.235.6, 21}
```

- A **half association** is either

- {protocol, local-address, local-process} or
- {protocol, foreign-address, foreign-process},
 - each describing one end of a connection (i.e., a socket)

SOCKET SYSTEM CALLS

The **socket** system call is used to **obtain a socket descriptor** representing an incompletely specified communication endpoint

```
#include <sys/types.h>
#include <sys/socket.h>
int socket (int family, int type, int protocol);
```

- For our purposes, we set
 - **family** to **AF_INET**
 - i.e., address family / Internet protocols
 - **type** to **SOCK_STREAM**
 - specifies sequenced, reliable, two-way, connection-oriented byte streams and is typically implemented with TCP
 - **protocol** to **0**
 - this is the default since the type **SOCK_STREAM** has only one protocol available, namely, **TCP**
- If successful, **socket** returns a non-negative integer corresponding to a **socket descriptor**
- If unsuccessful, **socket** returns **-1** and sets **errno**
- The **socket** system call fills in the **protocol** element of the association 5-tuple

SOCKET SYSTEM CALLS

Example

- `socket` system call
 - both client and server

```
int socketFD;  
socketFD = socket (AF_INET, SOCK_STREAM, 0);
```

andreeds.github.io

SOCKET SYSTEM CALLS

The **bind** system call is used by a **TCP server** to associate the handle for a socket communication endpoint with a logical network connection

```
#include <sys/types.h>
#include <sys/socket.h>
int bind (int sockfd, struct sockaddr *myaddr, int addrlen);
```

- **sockfd** is the **socket descriptor** returned by a previous call to **socket**
- **myaddr** is a **pointer** to a variable of type **sockaddr** (structure)
- **addrlen** is the **number of bytes** in a **sockaddr** structure
- If successful, **bind** returns **0**
- If unsuccessful, **bind** returns **-1** and sets **errno**
- The bind system call fills in the **local-address** and **local-process** elements of the association 5-tuple

SOCKET SYSTEM CALLS

Example

- **bind** system call
 - **server only**

```
int socketFD;  
struct sockaddr in serverAddress;
```

```
bind (socketFD, (struct sockaddr in *) &serverAddress, sizeof (serverAddress));
```

andreeds.github.io

SOCKET SYSTEM CALLS

The **listen** system call is used by a **TCP server** to indicate that it is ready to receive connection requests and to specify the number of waiting connection requests

```
#include <sys/socket.h>
int listen (int sockfd, int backlog);
```

- **backlog** specifies how many connection requests can be queued by the system while it waits for the server to return from handling a previous request
 - i.e., the server needs time to fork a child process and return
- If successful, **listen** returns **0**
- If unsuccessful, **listen** returns **-1** and sets **errno**
- When a socket is created by the socket **system** call, it is assumed to be an **active socket**
 - server and client
- An **active socket** is expected to eventually issue a connect system call
 - i.e., to set up the bidirectional communication with another socket
- The **listen** system call converts an unconnected socket into a **passive socket**
 - i.e., it should accept incoming connection requests

SOCKET SYSTEM CALLS

Example

- **listen** system call
 - **server only**

```
int socketFD;  
listen (socketFD, 1);
```

andreeds.github.io

SOCKET SYSTEM CALLS

The **accept** system call is used by a **TCP server to complete the logical connection between the server and the client**

```
#include <sys/types.h>
#include <sys/socket.h>
int accept (int sockfd, struct sockaddr *client, int *addrlen);
```

- The **client** and **addrlen** parameters are used to return the address of the connected process
- The **accept** system call is used by a TCP server to **return the next completed connection from the head of the connection queue established by the listen system call**
 - Automatically creates a new socket descriptor
 - After a connection request is received and accepted, a server typically forks, with the child process servicing the connection on the new socket descriptor and the parent process waiting for another connection request
 - If there are no connection requests pending, **accept** blocks
- If successful, **accept** returns a **new socket descriptor**
- If unsuccessful, **accept** returns **-1** and sets **errno**
- The **accept** system call can fill in the **foreign-address** and **foreign-process** elements of the association 5-tuple in the child process

SOCKET SYSTEM CALLS

Example

- **accept** system call
 - server only

```
int socketFD;  
int newSocketFD;  
struct sockaddr_in clientAddress;  
socklen_t addressLength;
```

```
newSocketFD = accept (socketFD, (struct sockaddr_in *) &clientAddress,  
&addressLength);
```

andredes.github.io

SOCKET SYSTEM CALLS

The **connect** system call is used by a TCP client to establish a link to the known port of the server

```
#include <sys/types.h>
#include <sys/socket.h>
int connect (int sockfd, struct sockaddr *servaddr, int addrlen);
```

- **sockfd** is the socket descriptor returned by a previous call to **socket**
- **servaddr** is a pointer to a variable of type **sockaddr** (structure)
- **addrlen** is the number of bytes in the **sockaddr** (structure)
- The client does not have to bind to a local address before a call to **connect**
- The **connect** system call fills in the **local-address**, **local-process**, **foreign-address**, **foreign-process** elements of the association 5-tuple
- If successful, **connect** returns **0**
- If unsuccessful, **connect** returns **-1** and sets **errno**

SOCKET SYSTEM CALLS

Example

- **connect** system call
 - **client only**

```
int socketFD;  
struct sockaddr in serverAddress;
```

```
connect (socketFD, (struct sockaddr *) &serverAddr, sizeof (serverAddr));
```

andreeds.github.io

OTHER USEFUL FUNCTIONS

Byte Manipulation

- There are multibyte fields in the various socket address structures that need to be manipulated
- These multibyte fields may not necessarily be standard C-strings (i.e., null terminated)

```
#include <string.h>
```

```
memcpy (char *dest, char *src, int bytes);
```

- `memcpy` moves the specified number of bytes from the source to the destination

```
memset (char *dest, int ch, int bytes);
```

- `memset` writes the character `ch` (as an **unsigned char**) into the specified number of bytes in the destination

```
int memcmp (char *ptr1, char *ptr2, int bytes);
```

- `memcmp` compares two arbitrary bytes strings, returning `0` if the strings are equal, a value greater than `0` if `*ptr1` is greater than `*ptr2`, and a value less than `0` if `*ptr1` is less than `*ptr2`

OTHER USEFUL FUNCTIONS

Byte Ordering

- There may be differences in byte order between different computer architectures and different network protocols
- The following four functions can be used to handle the conversions required.

```
#include <sys/types.h>
```

```
#include <netinet/in.h>
```

```
u_long htonl (u_long hostlong);    // host-to-network long  
u_short htons (u_short hostshort); // host-to-network short  
u_long ntohl (u_long netlong);     // network-to-host long  
u_short ntohs (u_short netshort);  // network-to-host short
```

OTHER USEFUL FUNCTIONS

Address Conversion

- An Internet address is usually written in dotted-decimal format (e.g., 192.43.235.1)
- The following two functions are used to convert between dotted-decimal format and that contained in an `in_addr` structure.

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
```

```
unsigned long inet_addr (char *ptr);
```

- `inet_addr` converts a character string in dotted-decimal format to a 32 bit Internet address

```
char *inet_ntoa (struct in_addr inaddr);
```

- `inet_ntoa` converts a 32 bit Internet address to dotted-decimal format

OTHER USEFUL FUNCTIONS

- We usually refer to computers by their human-readable names
- However, computers refer to computers by their numerical addresses
- Consequently, library routines are provided to convert between human-readable names and numeric addresses

andreeds.github.io

OTHER USEFUL FUNCTIONS

```
#include <unistd.h>
int gethostname (char *hostname, int len);
```

- The returned parameter, `gethostname`, is null-terminated
- If successful, `gethostname` returns `0`
- If unsuccessful, `gethostname` returns `-1`

andreds.github.io

OTHER USEFUL FUNCTIONS

```
#include <netdb.h>
```

```
struct hostent *gethostbyname (char *hostname);
```

- The `gethostbyname` function returns a pointer to a `hostent` structure.

```
struct hostent
{
    char *h_name;           // official name of host
    char **h_aliases;
    int  h_addrtype;        // host address type (always contains AF_INET)
    int  h_length;          // length of address (always contains 4)
    char **h_addr_list;     // null terminated list of addresses from name server
};

#define h_addr h_addr_list [0]
```

- The `h_addrtype` field always contains `AF_INET`
- The `h_length` field always contains 4 (i.e., the length of an Internet address in bytes)
- The `h_addr_list` field is actually an array of pointers to structures of type `in_addr`

PIPE SYSTEM CALLS

1ServerExampleReadFile.c • 1ClientExampleReadFile.c

URCourses

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
```


PIPE SYSTEM CALLS

2socket[Client/Server/Functions/Include]Hilderman.[c/h]

URCourses

```
#include "socketInclude.h"
#include "socketFunctions.h"

int main (int argc, char **argv)
{
    int i;
    int socket_fd;
    struct sockaddr_in server_addr;
    struct hostent *hp;

    if (argc < 2)
    {
```

PIPE SYSTEM CALLS

3TCP[Client/Server].c

URCourses

```
#include <stdio.h>
#include <netdb.h>
#include <netinet/in.h>
#include <stdlib.h>
#include <string.h>
#include <sys/socket.h>
#include <sys/types.h>
#define MAX 80
#define PORT 8080
#define SA struct sockaddr

// Function designed for chat between
client and server.
void func(int sockfd)
```