

Programming Signals

A. A Review of Interrupts:

1. A *hardware interrupt*:

- a. A signal sent to the processor from a device that is part of the computer itself or from some other external device.
- b. Causes the CPU to stop its current task and transfer execution to the kernel.
- c. The act of initiating a hardware interrupt is referred to as an *interrupt request* (a.k.a. IRQ).
- d. An IRQ has a *type* associated with it that is determined by where and/or how it originated.
- e. Example

The disk controller indicating that a read or write operation has completed.

- f. Example – External device

Clicking the mouse or entering a character at the keyboard.

- g. The kernel *services* the IRQ using a table of pointers in memory, where each pointer represents the start address of the *interrupt service routines* associated with each type.

2. A *software interrupt*:

- a. Caused by an exceptional condition in the CPU (a.k.a. *trap* or *exception*) or a special instruction in the instruction set that causes an interrupt when it is executed.
- b. Could be an error or event during execution that is exceptional or serious enough that it cannot be handled within the currently executing task.
- c. Causes the CPU to stop its current task and transfer execution to the kernel.
- d. Example

If the CPU's arithmetic logic unit attempts to divide a number by zero.

- e. Example

If an attempt is made to reference memory with an uninitialized pointer.

- f. Could be a system call to request services from low-level system software, such as device drivers.

- g. Example

An `open`, `read`, or `write` instruction results in a *system call* that causes the CPU to stop its current task, trap to the *system call interface*, and transfer execution to the kernel.

- h. Example

A signal.

- i. Signals can occur in the middle of executing an instruction in an application program, so care is required when using them (i.e., such as whether the interrupted code is *re-entrant*).

- Re-entrant functions are those that can be called more than once by the same process or simultaneously by multiple processes.
- A re-entrant function is written so that none of its code is modifiable (no values are changed) and it does not keep track of anything.
- The calling programs keep track of their own progress (variables, flags, etc.), thus one copy of the re-entrant code can be shared by any number of users or processes.

j. Functions that are re-entrant do not:

- Use static internal data structures (which could be modified by the signal handler).
- Call `malloc` or `free` (which stores allocated areas in a linked list that could be corrupted by the signal handler).
- Use the standard I/O library (the standard I/O library uses some global data structures which could be modified by the signal handler).

B. Signals

1. A *signal* is software interrupt mechanism that generates a notification indicating to a process that some event has occurred.
2. Every signal has a name and is associated with an integer-valued number.
3. Example – Partial list of Linux signals (could be different than UNIX, other Linuxes, or other OSs as values are system-dependent (can be found in `/usr/include/bits/signum.h` on Linux machines in CL115)).

```
#define SIGHUP      1      /* hangup (POSIX) */
```

```

#define SIGINT      2      /* terminal interrupt (ANSI) */
#define SIGQUIT     3      /* terminal quit (POSIX) */
#define SIGILL      4      /* illegal instruction (ANSI) */
#define SIGTRAP     5      /* trace trap (POSIX) */
#define SIGABRT     6      /* abort (4.2 BSD) */
#define SIGBUS      7      /* bus error (4.2 BSD) */
#define SIGFPE      8      /* floating point exception (ANSI) */
#define SIGKILL     9      /* kill (can't be caught or ignored) (POSIX) */
#define SIGUSR1    10      /* user defined signal 1 (POSIX) */
#define SIGSEGV    11      /* segmentation violation (ANSI) */
#define SIGUSR2    12      /* user defined signal 2 (POSIX) */
#define SIGPIPE    13      /* write on a pipe with no reader (POSIX) */
#define SIGALRM    14      /* alarm clock (POSIX) */
#define SIGTERM    15      /* termination signal from kill (ANSI) */
#define SIGSTKFLT  16      /* stack fault */
#define SIGCHLD    17      /* child status change */
#define SIGCONT    18      /* if stopped, continue executing (POSIX) */
#define SIGSTOP    19      /* stop (can't be caught or ignored) (POSIX) */
#define SIGTSTP    20      /* terminal stop (POSIX) */
#define SIGTTIN    21      /* background process trying to read from terminal (POSIX)
    */
#define SIGTTOU    22      /* background process trying to write to terminal (POSIX)
    */
#define SIGURG     23      /* urgent condition related to socket (4.2 BSD) */
#define SIGXCPU    24      /* cpu limit exceeded (4.2 BSD) */
#define SIGXFSZ    25      /* file size limit exceeded (4.2 BSD) */
#define SIGVTALRM  26      /* virtual alarm clock (4.2 BSD) */
#define SIGPROF    27      /* profiling alarm clock (4.2 BSD) */
#define SIGWINCH   28      /* window size change (4.3 BSD) */
#define SIGIO      29      /* I/O now possible (4.2 BSD) */

```

```
#define SIGPWR    30           /* power failure restart (System V) */
```

4. Several events can cause a signal to be generated:

- a. *Hardware exceptions*: The conditions are detected by the hardware, which notifies the kernel, which generates the appropriate signal, which is sent to the appropriate process. Examples include:
 - Division by zero (i.e., SIGFPE).
 - Invalid memory reference (i.e., SIGSEGV).
- b. *Software conditions*: When an event happens that a process should know about. Examples include:
 - Writing to a pipe that has no reader (i.e., SIGPIPE).
 - When a timer set by a process expires (i.e., SIGALRM).
 - When some user-defined condition occurs (i.e., SIGUSR1).
- c. *Terminal-generated signals*: When a user presses keys simultaneously in particular combinations. Examples include:
 - Control/C to stop a runaway process (i.e., SIGINT).
 - Control/Z to suspend a process running in foreground (i.e., SIGTSTP).
- d. *The kill/sigqueue system calls* (more on these later): To send any signal from a user-owned process to any other user-owned process.

- e. *The kill command*: This is a command line interface to the kill system call to enable a signal to be sent from the shell to a (typically) runaway background process.
5. There are *two* generations of signals (at least for the purposes of our discussion there is):
- a. *Unreliable*: A throwback to the very early versions of signals in UNIX that have been superseded by the POSIX signals standard.
 - b. *Reliable*: A (modern) version of signals adhering to the POSIX signals standard.
6. Unreliable signals suffer from a number of problems and should not be used in new programs:
- a. They can get lost (i.e., a signal could be sent but the intended recipient misses it).
 - b. The disposition of a signal set by a process must be reset by the process each time the signal is received.
 - If the disposition is to catch the signal (with a *signal handler*), but the default action is to kill the process, there is a small window of time where the default action would be enabled until the process resets it again.
 - Another example of a *race condition*.
 - c. They handling of a signal cannot be deferred, only ignored.
7. Reliable signals solve the problems with unreliable signals.
- a. The disposition of a signal set by a process is not reset to the default each time a signal is received, only when the process specifically changes it.
 - b. Processes have the ability to both ignore or temporarily block signals.

- When a signal is blocked by a process, the kernel places it on a queue of pending signals for that process.
- A blocked signal remains pending until the process unblocks it or changes its disposition to ignore it.
- SIGKILL and SIGSTOP cannot be blocked.

8. From here on, we assume the use of reliable signals.

9. A signal will always be in one of *three* possible states:

- a. A signal is *generated* (i.e., sent to a process) when the event that causes the signal occurs.
- b. A signal is *pending* (i.e., blocked) if it has been generated but not delivered.
- c. A signal is *delivered* when the action associated with the signal is actually invoked.

10. The *lifetime* of a signal is the interval between its generation and delivery.

11. Signals may be generated in *two* ways:

- a. *Synchronously*: When an event occurs that is directly caused by the execution of a process' code (also called a *trap*) (e.g., SIGFPE).
- b. *Asynchronously*: When an event occurs at a seemingly random time with respect to the process (e.g., SIGKILL).

12. A process can respond to the receipt of a signal (called the signal's *disposition* or *associated action*) in *two* ways when it is delivered:

- a. *Catch it*: Call a signal handler, a user-written function contained in a process that describes how the event should be handled. Examples include:
 - Catching `SIGTERM` (the default termination signal sent by the `kill` command) to release memory and delete temporary files.
 - Catching `SIGCHLD` to catch the termination of a child process.
- b. *Take one of five possible default actions*:
 - Ignore the signal (it is possible to ignore certain signals generated by a hardware exception (e.g., `SIGFPE`), but process behaviour may become difficult to understand).
 - Terminate the process.
 - Core dump.
 - Stop if the process is currently running.
 - Continue if the process is currently stopped.

13. Generate signals from the shell with the `kill` command.

- a. List the symbolic names of the signals available (POSIX).
- b. Kill a particular process (POSIX).

```
kill -l
```



```
kill -s signal_name pid
```

c. Example

```
kill -s USR1 3423
```

d. Traditional `kill` command (still supported by POSIX, but only because of widespread usage).

```
kill -signal_name pid
```

```
kill -signal_number pid
```

e. Example

```
kill -KILL 3423
```

```
kill -9 3423
```

14. The `kill` system call is used to send a signal to a process.

```
#include <sys/types.h>
```

```
#include <signal.h>
```

```
int kill (pid_t pid, int sig);
```

- a. The `kill` system call sends the signal specified by *sig* to the process specified by *pid*.
- b. The first parameter, *pid*, is a valid process identifier (can actually have other values, but we don't discuss them here).
- c. The second parameter, *sig*, must be a valid signal name or 0.

- If `sig` is 0, (i.e., the NULL signal), normal error checking is performed, *but no signal is actually sent*.
 - Why would we want to do this? We can use 0 to check whether *pid* is a valid process before we actually try to kill it.
- d. If successful, `kill` returns 0. If unsuccessful, `kill` returns -1 and sets `errno`.
- e. Example – Child killing its parent.

PROGRAM = `childKillsParent.c`

15. The POSIX signal handling interface makes use of *signal sets* rather than individual signals.

- a. A signal set is a *bit mask*, one bit for each signal.
- b. If a bit is set to 0 (1), the corresponding signal is not (is) a member of the set.
- c. Since the number of different signals can exceed the number of bits in an `int`, a signal set is of the `sigset_t` data type (defined in `signal.h`).

16. Signal sets can be created and deleted using the following *five* functions:

```
#include <signal.h>

int sigemptyset (sigset_t *signal_set) ;
int sigfillset (sigset_t *signal_set) ; (Won't be discussed further in this course)
int sigaddset (sigset_t *signal_set, int signal_number) ;
int sigdelset (sigset_t *signal_set, int signal_number) ; (Won't be discussed further in this course)
```

All four return 0 on success. Otherwise, they return -1.

```
int sigismember (sigset_t *signal_set, int signal_number) ;
```

 (Won't be discussed further in this course)

It returns 1 if true. Otherwise, it returns 0.

17. The `sigemptyset` function initializes the signal set pointed to by *signal_set* to exclude all signals (i.e., the empty set).

18. Example

```
sigset_t interruptMask;  
  
sigemptyset (&interruptMask);
```

19. The `sigaddset` function adds the single signal specified by *signal_number* to the signal set pointed to by *signal_set*.

20. Example

```
sigaddset (&interruptMask, SIGRTMIN);
```

21. The action associated with a signal in a signal set can be examined and modified using the following *two* functions:

```
int sigaction (int signal_number, const struct sigaction *action, struct sigaction *old_action) ;  
sigprocmask (int what_to_do, const sigset_t *signal_set, sigset_t *old_signal_set) ;
```

22. The `sigaction` function allows for the examination and modification of the action associated with a particular signal.

- a. A sigaction struct consists of *four* members:

```
struct sigaction
{
    int sa_flags;
    void (*sa_handler) ();
    void (*sa_sigaction) (int, siginfo_t *, void *);
    sigset_t sa_mask;
};
```

- b. If sa_flags is set to SA_SIGINFO, then *three* arguments can be passed to the signal handler if its function prototype is declared as follows (enables the *receiving* process to determine the identity of the *sending* process):

```
static void signal_handler (int signal_number, siginfo_t *signal_info, void *context);
```

- c. The siginfo_t data type provides information about why a signal was generated and where it originated.

```
typedef struct
{
    int si_signo;
    int si_errno;
    int si_code;
    pid_t si_pid;
    uid_t si_uid;
    .
    .
    .
} siginfo_t;
```

23. Example

```
struct sigaction act;  
.  
.  
.  
act.sa_sigaction = &SignalHandler;  
act.sa_flags = SA_SIGINFO;  
sigemptyset (&act.sa_mask);  
sigaction (SIGRTMIN, &act, NULL);  
.  
.  
.
```

24. The `sigprocmask` function allows for the examination and modification of the signal mask stored in the `sigaction` struct.

25. Example

```
.  
.  
.  
sigprocmask (SIG_BLOCK, &interruptMask, NULL);  
.  
.  
.  
do something you don't want interrupted  
.  
.  
.
```

```
sigprocmask (SIG_UNBLOCK, &interruptMask, NULL);  
.  
.  
.
```

26. The `sigqueue` function is an extension to the `kill` function that put signals in a queue.

```
int sigqueue (pid_t pid, int signal_number, const union sigval value);
```

- a. Returns 0 on success. Otherwise, it returns -1.

27. Example

```
union sigval dummyValue;  
  
sigqueue (pid, SIGRTMIN, dummyValue);
```

28. The `pause` system call suspends a process until a signal is delivered.

```
#include <unistd.h>  
  
int pause (void);
```

- a. The `pause` system call suspends the calling process until it receives a signal that it is not currently ignoring.
- b. The `pause` system call returns -1 and sets `errno` when a signal handler is executed.
- c. If the disposition of the received signal is to terminate, `pause` does not return. Otherwise, the process continues executing from where it was suspended.

29. Example – A process catching a signal sent from the command line.

PROGRAM = `catchSignals.c`

30. Example – A process catching a signal sent from a child.

PROGRAM = `parentCatchSignals.c`