

UNIVERSITY OF REGINA

CS330-001 INTRODUCTION TO OPERATING SYSTEMS

ANDRÉ E. DOS SANTOS

dossantos@cs.uregina.ca

andreeads.github.io

CS330-001
INTRODUCTION TO
OPERATING SYSTEMS

PROGRAMMING SIGNals C

ANDRÉ E. DOS SANTOS
dossantos@cs.uregina.ca
andreeds.github.io

WHAT ARE SIGNALS

- A **signal** is software interrupt mechanism that generates a notification indicating to a process that some event has occurred
- Every signal has a name and is associated with an integer-valued number

SIGNALS

```
#define SIGHUP      1      /* hangup (POSIX) */
#define SIGINT      2      /* terminal interrupt (ANSI) */
#define SIGQUIT     3      /* terminal quit (POSIX) */
#define SIGILL      4      /* illegal instruction (ANSI) */
#define SIGTRAP     5      /* trace trap (POSIX) */
#define SIGABRT     6      /* abort (4.2 BSD) */
#define SIGBUS      7      /* bus error (4.2 BSD) */
#define SIGFPE      8      /* floating point exception (ANSI) */
#define SIGKILL     9      /* kill (can't be caught or ignored) (POSIX) */
#define SIGUSR1    10      /* user defined signal 1 (POSIX) */
#define SIGSEGV    11      /* segmentation violation (ANSI) */
#define SIGUSR2    12      /* user defined signal 2 (POSIX) */
#define SIGPIPE    13      /* write on a pipe with no reader (POSIX) */
#define SIGALRM    14      /* alarm clock (POSIX) */
#define SIGTERM    15      /* termination signal from kill (ANSI) */
#define SIGSTKFLT  16      /* stack fault */
#define SIGCHLD    17      /* child status change */
#define SIGCONT    18      /* if stopped, continue executing (POSIX) */
#define SIGSTOP    19      /* stop (can't be caught or ignored) (POSIX) */
#define SIGTSTP    20      /* terminal stop (POSIX) */
#define SIGTTIN    21      /* background process trying to read from terminal (POSIX) */
#define SIGTTOU    22      /* background process trying to write to terminal (POSIX) */
#define SIGURG     23      /* urgent condition related to socket (4.2 BSD) */
#define SIGXCPU    24      /* cpu limit exceeded (4.2 BSD) */
#define SIGXFSZ    25      /* file size limit exceeded (4.2 BSD) */
#define SIGVTALRM  26      /* virtual alarm clock (4.2 BSD) */
#define SIGPROF    27      /* profiling alarm clock (4.2 BSD) */
#define SIGWINCH   28      /* window size change (4.3 BSD) */
#define SIGIO      29      /* I/O now possible (4.2 BSD) */
#define SIGPWR     30      /* power failure restart (System V) */
```

GENERATING SIGNALS

Hardware exceptions

- The conditions are detected by the hardware, which notifies the kernel, which generates the appropriate signal, which is sent to the appropriate process. Examples include:
 - Division by zero (i.e., **SIGFPE**).
 - Invalid memory reference (i.e., **SIGSEGV**).

Software conditions

- When an event happens that a process should know about. Examples include:
 - Writing to a pipe that has no reader (i.e., **SIGPIPE**).
 - When a timer set by a process expires (i.e., **SIGALRM**).
 - When some user-defined condition occurs (i.e., **SIGUSR1**).

Terminal-generated signals

- When a user presses keys simultaneously in particular combinations. Examples include:
 - Control/C to stop a runaway process (i.e., **SIGINT**).
 - Control/Z to suspend a process running in foreground (i.e., **SIGTSTP**).

GENERATING SIGNALS

- There are two generations of signals (at least for the purposes of our discussion there is):

Unreliable

- A throwback to the very early versions of signals in UNIX that have been superseded by the POSIX signals standard.

Reliable

- A (modern) version of signals adhering to the POSIX signals standard.

UNRELIABLE SIGNALS

- Unreliable signals suffer from a number of problems and should not be used in new programs:
 - They can get lost (i.e., a signal could be sent but the intended recipient misses it)
 - The disposition of a signal set by a process must be reset by the process each time the signal is received
 - If the disposition is to catch the signal (with a signal handler), but the default action is to kill the process, there is a small window of time where the default action would be enabled until the process resets it again
 - Another example of a race condition
 - The handling of a signal cannot be deferred, only ignored.

RELIABLE SIGNALS

- Reliable signals solve the problems with unreliable signals
 - The disposition of a signal set by a process is not reset to the default each time a signal is received, only when the process specifically changes it
 - Processes have the ability to both ignore or temporarily block signals
 - When a signal is blocked by a process, the kernel places it on a queue of pending signals for that process
 - A blocked signal remains pending until the process unblocks it or changes its disposition to ignore it
 - **SIGKILL** and **SIGSTOP** cannot be blocked
- From here on, we assume the use of reliable signals.

SIGNAL STATES

- A signal will always be in one of three possible states
 1. A signal is **generated** (i.e., sent to a process) when the event that causes the signal occurs.
 2. A signal is **pending** (i.e., blocked) if it has been generated but not delivered.
 3. A signal is **delivered** when the action associated with the signal is actually invoked
- The **lifetime** of a signal is the interval between its generation and delivery.

SIGNAL GENERATION

- Signals may be generated in two ways:
 1. **Synchronously**: When an event occurs that is directly caused by the execution of a process' code (also called a trap) (e.g., `SIGFPE`)
 2. **Asynchronously**: When an event occurs at a seemingly random time with respect to the process (e.g., `SIGKILL`)

SIGNAL RESPONSE

A process can respond to the receipt of a signal (called the **signal's disposition** or **associated action**) in two ways when it is delivered:

- **Catch it:** Call a signal handler, a user-written function contained in a process that describes how the event should be handled
 - Examples include:
 - Catching **SIGTERM** (the default termination signal sent by the **kill** command) to release memory and delete temporary files
 - Catching **SIGCHLD** to catch the termination of a child process
- **Take one of five possible default actions:**
 - Ignore the signal
 - Terminate the process
 - Core dump
 - Stop if the process is currently running
 - Continue if the process is currently stopped

kill COMMAND

- List the symbolic names of the signals available (POSIX)

```
kill -l
```

- Kill a **particular** process (POSIX)

```
kill -s signal_name pid
```

- Traditional kill command (still supported by POSIX, but only because of widespread usage)

```
kill -signal_name pid
```

```
kill -signal_number pid
```

ELEMENTARY SIGNAL SYSTEM CALLS

The **kill** system call is used to send a signal to a process

```
#include <sys/types.h>
#include <signal.h>
int kill (pid_t pid, int sig);
```

- The **kill** system call sends the signal specified by **sig** to the process specified by **pid**
- **pid** is a valid process identifier
- **sig** must be a valid signal name or **0**
 - If sig is 0, (i.e., the NULL signal), normal error checking is performed, but no signal is actually sent
 - Why would we want to do this?
 - We can use 0 to check whether pid is a valid process before we actually try to kill it
- If **successful**, kill returns **0**
- If **unsuccessful**, kill returns **-1** and sets **errno**

ELEMENTARY SIGNAL SYSTEM CALLS

childKillsParent.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>
#include <signal.h>
```

```
int main ()
{
    pid_t childPid;
    int status;
    pid_t waitPid;
```

ELEMENTARY SIGNAL SYSTEM CALLS

catchSignals.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>
#include <signal.h>
```

```
int main ()
{
    pid_t childPid;
    int status;
    pid_t waitPid;
```

ELEMENTARY SIGNAL SYSTEM CALLS

parentCatchSignals.c

URCourses

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>
#include <signal.h>

int main ()
{
    pid_t childPid;
    int status;
    pid_t waitPid;
```